

Sluice: Network-Wide Data Plane Programming

Vikas Natesh¹ Pravein Govindan Kannan² Anirudh Sivaraman¹ Ravi Netravali³
¹New York University ²National University of Singapore ³University of California, Los Angeles

1 INTRODUCTION

The last several years have seen the emergence of programmable network devices including both programmable switching chips and programmable network interface cards (NICs). Along with the rise of x86-based packet processing for middleboxes and virtual switches, these trends point towards a future where the entire network will be programmable. The benefits of network programmability range from commercial use cases such as network virtualization implemented on the programmable Open vSwitch platform to more recent projects that implement packet scheduling, measurement, and application offload of niche applications on programmable switches.

Benefits are clear, but are difficult to reap as programming network as a whole remains challenging. Current programming languages target individual network devices, e.g., P4 for the Tofino programmable switching chip and the Netronome programmable NIC. However, at present, there is no unified programming model to express and implement general functionality at the level of an entire network, without having to individually program each network device.

Maple was an early example of a network-wide programming model designed for OpenFlow switches. Maple automatically divided functionality between a stateless component running on switches and a stateful component running on the network's controller. SNAP is a more recent example of network-wide programming; unlike Maple, it additionally offloads stateful functionality to switches by leveraging stateful processing available in programmable switches. However, both Maple and SNAP cannot express programmable-switch functionality that affects network performance at fine time scales RN: *is there a short 'why' for this?*, e.g., packet scheduling, congestion control, fine-grained measurement of microbursts, and load balancing. In other words, Maple and SNAP generate highly optimized code, but are restricted in what they can express.

This demo presents **Sluice**, a programming model that takes a high-level specification of a network program and compiles it into runnable code that can be launched directly on the programmable devices of a network. In contrast to prior network-wide programming models like SNAP and Maple, Sluice aims to be more general RN: *should directly contrast with Maple/SNAP being coarse-grained*, but potentially at the cost of quality of generated code. Sluice endows network operators with the ability to design and deploy large




Figure 1: Sluice Workflow

network programs for various functions such as scheduling, measurement, and application offloading. The benefits of Sluice can be summarized as follows:

- Sluice provides the same functionality as P4 but makes it easier to program the data plane of an entire network.
- Sluice automatically reduces the amount of boilerplate code needed to write an equivalent program in P4. For instance, the 9 line traffic matrix sluice program translates into over 200 lines of P4 code.

We demonstrate Sluice's functionality and ease of use via two examples: traffic matrix generation for network analysis and a streaming join-filter operation.

2 SLUICE DESIGN

In the Sluice model, a network-wide program consists of high-level code *snippets* annotated by the operator to run on particular devices in a network. The code in each snippet is to be executed on packets arriving at its corresponding device. Snippets support a variety of operations: read-from/write-to packets; arithmetic using packet data, local variables, or stateful register arrays; and control flow statements. To handle computation on custom packet headers not supported by default (ip/tcp/udp/eth), users may define packet declarations similar to C structs. An optional annotation in the declaration, the parser condition, automatically generates a header parser. This lets the user restrict snippets to operate on specific flows or IP address ranges. Sluice programs may also import device-specific variables/attributes for use in code snippets.

Figure 1 describes the Sluice workflow. The compiler translates each snippet of a sluice program into a device-specific program. After initial parsing, lines of code in the snippet are decomposed into a directed acyclic graph (DAG) that maps dependencies between variables in each snippet RN: *I'm confused. Isn't each node a snippet? Or are you breaking a snippet into smaller components?*. This graph is then passed to the backend of the compiler that generates the corresponding P4 program for that device, for example bmv2 or Tofino.¹

¹Currently we only support bmv2 but plan to support more devices.




Figure 2: Topology For Traffic Matrix Demo

3 DEMONSTRATIONS

3.1 Traffic Matrix

Figure 2 displays the Mininet network topology for our traffic matrix demo. Packets are sent over UDP from each host to all other hosts according to a poisson traffic model with mean inter-arrival time of 0.5 seconds. The codelet below is our Sluice program with a single snippet *traffic_example* that is launched on all switches of the network. To run the simulation, the user passes the Sluice program and network topology to the compiler. The compiler generates P4 code to run on each switch as well as control plane table entries for routing packets through the topology.

```
import device psa;
```

```
packet p: udp(srcPort:1234)
  nhops : bit <32>;
```

```
@ bmv2 : ;
snippet traffic_example()
  persistent cnt : bit <32>[10];
  cnt[psa.ingress_port] = cnt[psa.ingress_port] + 1;
  p.nhops = p.nhops + 1;
```

This demo shows how a simple Sluice program can be used to enable each switch to measure link usage for a specific flow of user-defined packets on UDP srcPort 1234. Each packet *p* contains a custom header *nhops* that is incremented each time the packet enters a switch to inform the receiving host of the number of hops the packet took. Each switch maintains a stateful register counter *cnt*, indexed by switch ingress port, that tracks how many packets have entered through that ingress port. Aggregated over all switches, these counters data represent a matrix measuring each link's usage in the network at a given time. This matrix (residing on the whole network) is then queried once every second from the control plane to generate time-series plots of link-utilization for each link. Figure 3 displays the histogram of packet rates on link s1-s3 after collecting data for 3 minutes. The expected distribution of packet rates *poisson*($\mu = 2$ packets/sec) is also plotted to confirm the accuracy of the Mininet emulation.

3.2 Stream processing

This example demonstrates a simple join-filter operation between two streams of tuples. A stream is an unbounded table where a packet represents a tuple of data (*ad_id*, *impression_time*, *click_time*) enclosed in a custom header. The topology in Figure 4 describes the data flow and shows how




Figure 3: CDF of Packet Rate on link s1-s3

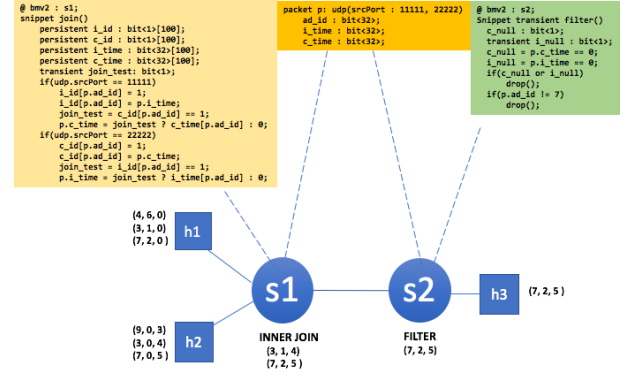


Figure 4: streaming example topology, data flow, and code placement on switches

an operator query runs on the switches of the network. Host 1 sends a stream of ad impressions while Host 2 sends a stream of ad clicks. The two streams are joined on the *ad_id* field at s1 and filtered on the *ad_id* field at s2 and the result is sent to h3.

4 FUTURE WORK

An optimizing Sluice compiler. We envision using the dependency DAG to provide several automatic optimizations or code transformations. For example, it is possible that certain lines of code in a snippet cannot be run on the device annotated by the operator, e.g., programmable switching chips have limited support for floating point or complex string operations. Code containing such features must be moved to the control plane or an end host data plane while at the same time, preserving the original program semantics intended by the operator. Doing this automatically would free the Sluice programmer from reasoning about these semantics.

Supporting multi-tenancy. Another area of future work is allowing Sluice RN: *would this really be a change to Sluice? Or would this be using Sluice to do something else?* to support multiple tenants with their own Sluice programs running on their own virtual networks overlayed on the same physical topology. If each tenant wants to run their own

~~Sluice: Network-Wide Data Plane Programming~~
network-wide program on their virtual topology, the network operator will need to merge all these into one data plane implementation that runs on the entire physical network. Extending Sluice to support this multi-tenancy use case would allow us to provide the same benefits to the data

plane that multi-tenant network virtualization [?] provided for the control plane.

5 REFERENCES